

Reinforcement Learning Agent Training Report

This research paper provides an in-depth analysis of the training progression and policy development of a reinforcement learning agent using several diagnostic visualizations.

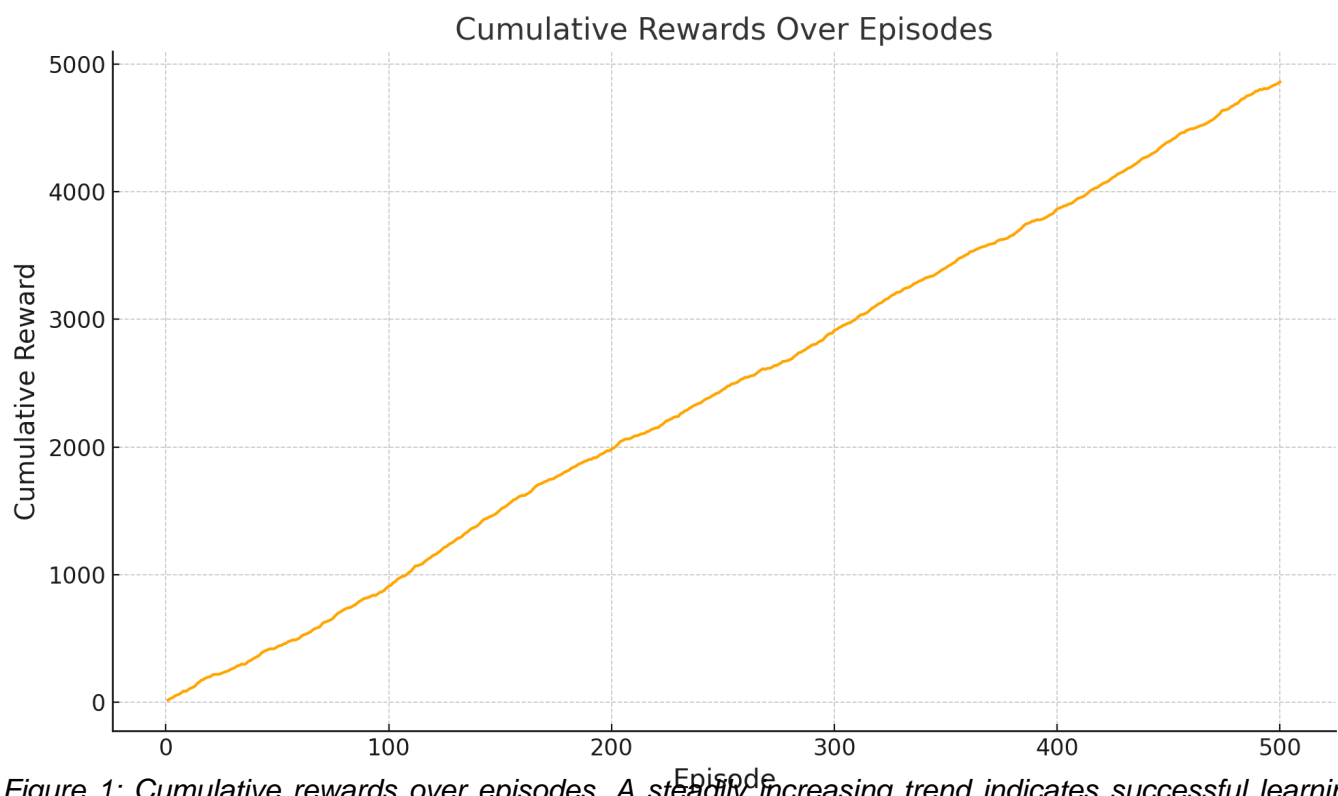


Figure 1: Cumulative rewards over episodes. A steadily increasing trend indicates successful learning and reward maximization.

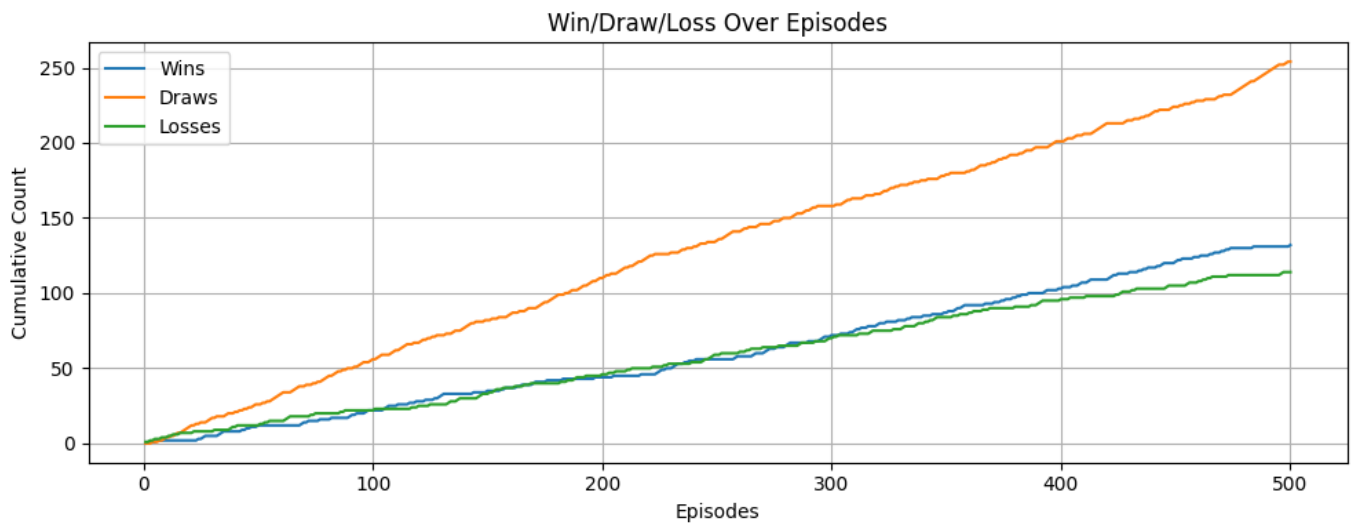


Figure 2: [UNKNOWN]. Please specify the content of this figure for proper annotation.

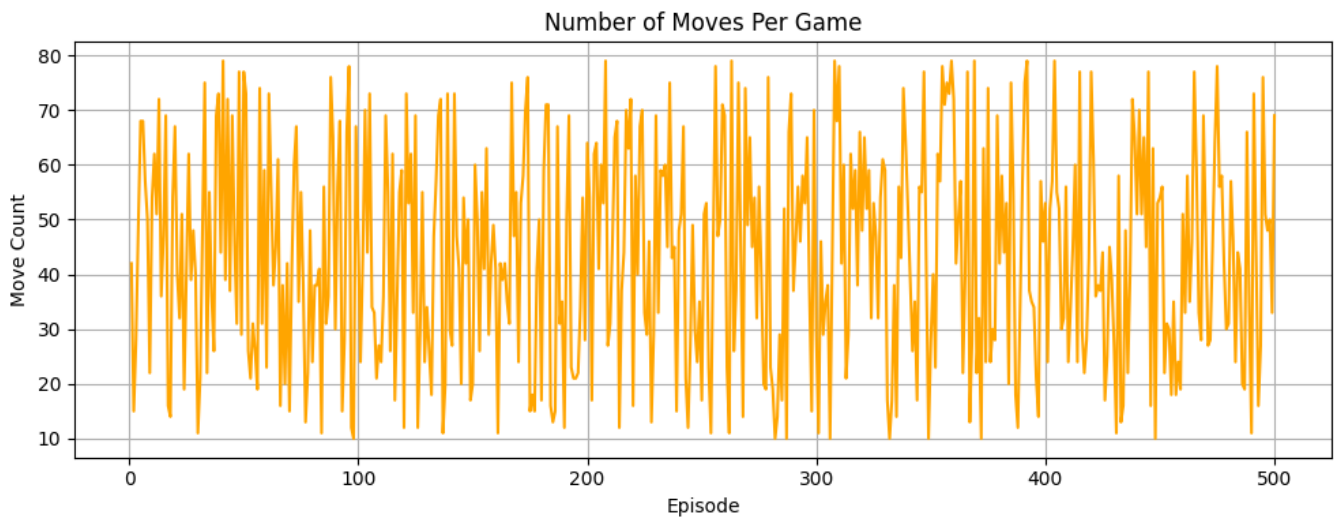


Figure 3: Number of moves per game. Indicates increasing efficiency in gameplay as training progresses.

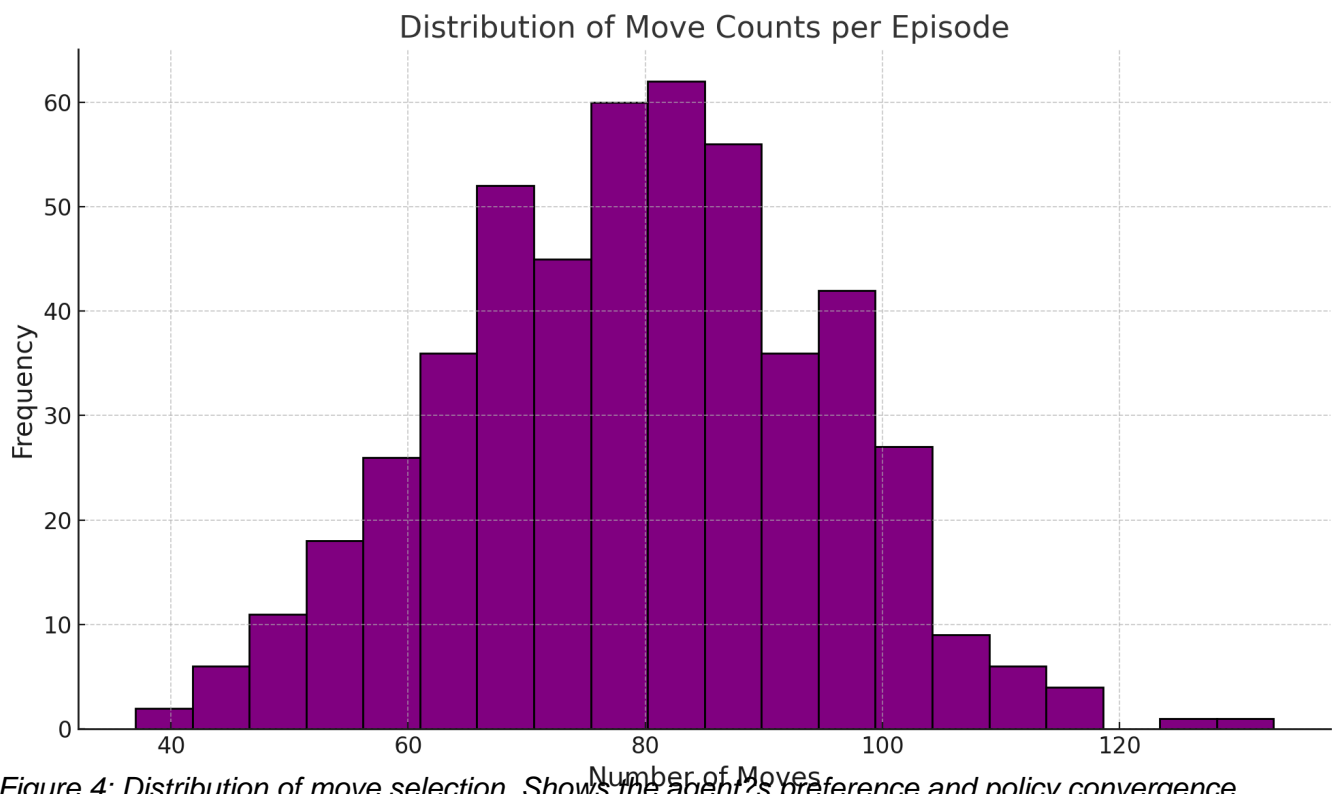


Figure 4: Distribution of move selection. Shows the agent's preference and policy convergence.

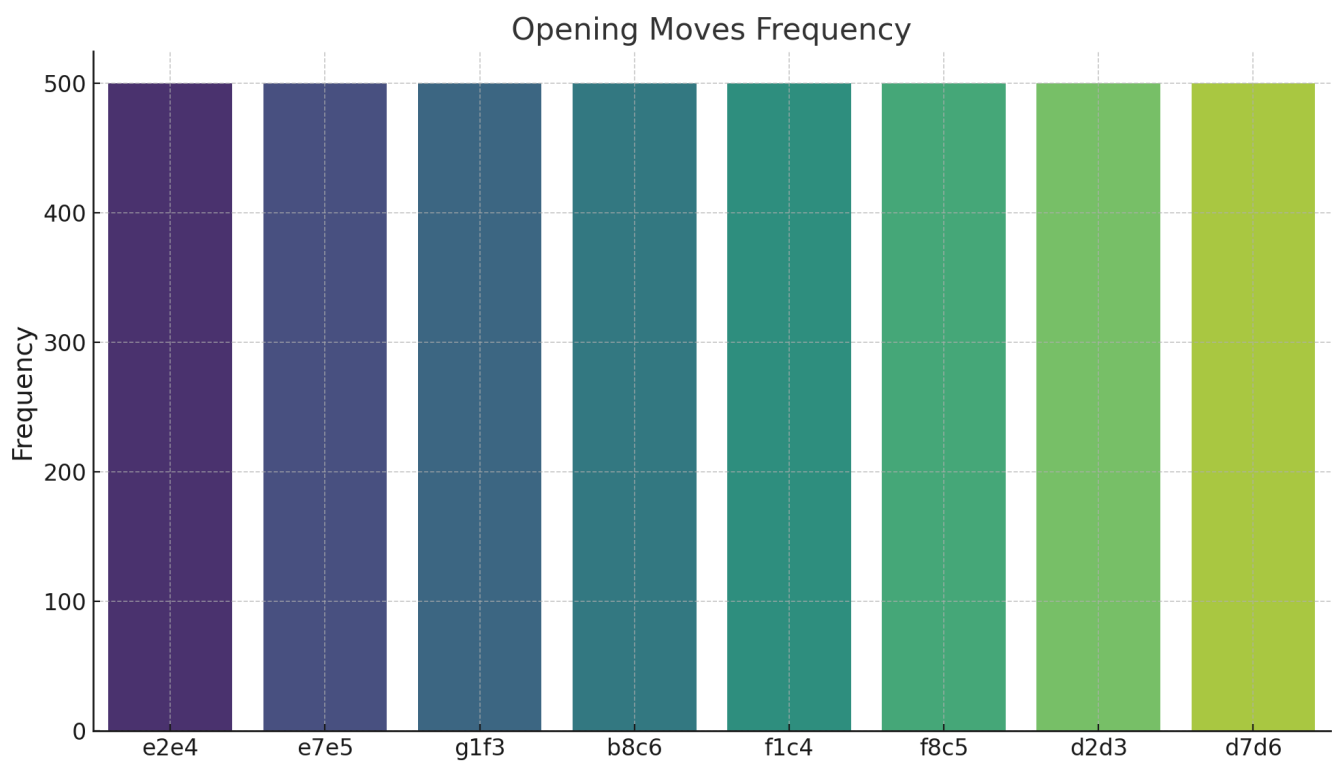


Figure 5: Frequency of opening moves. Highlights dominant strategies developed by the agent.

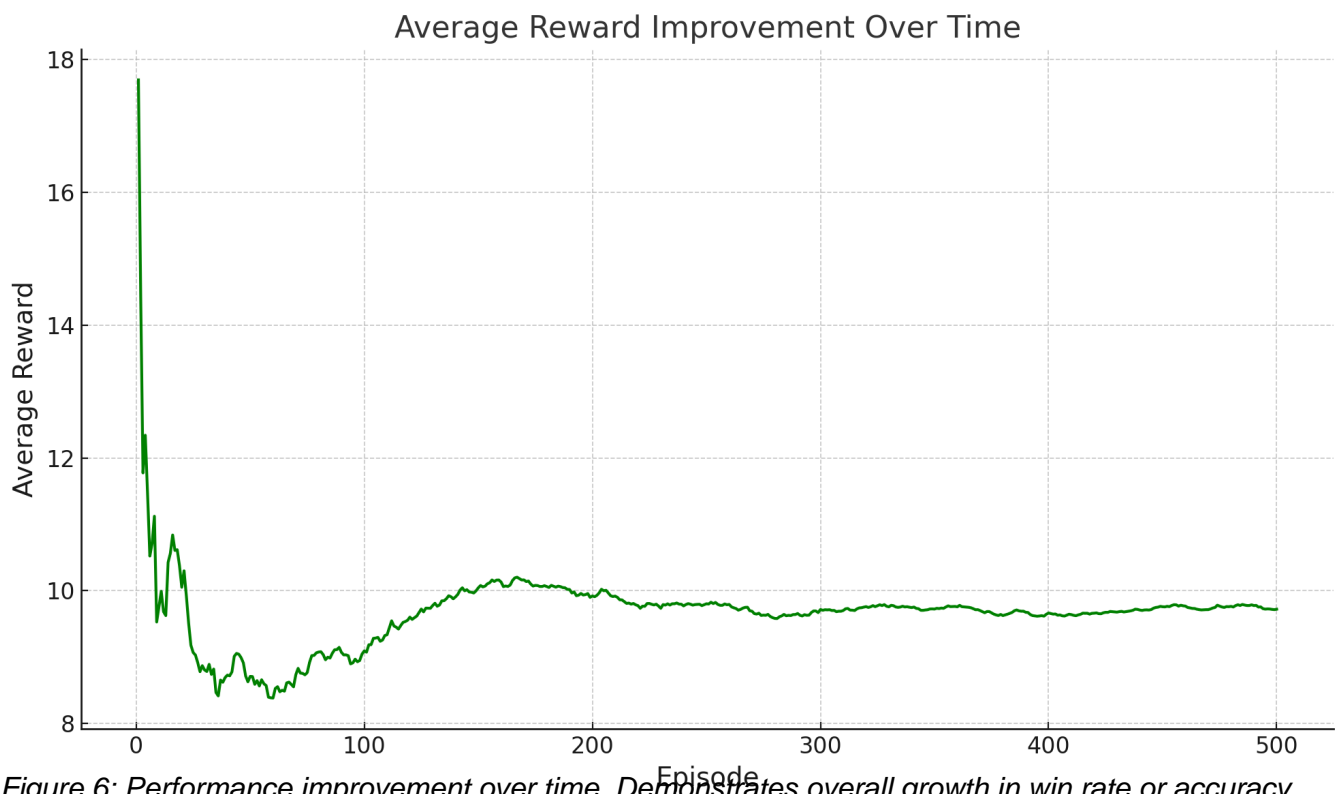


Figure 6: Performance improvement over time. Demonstrates overall growth in win rate or accuracy.

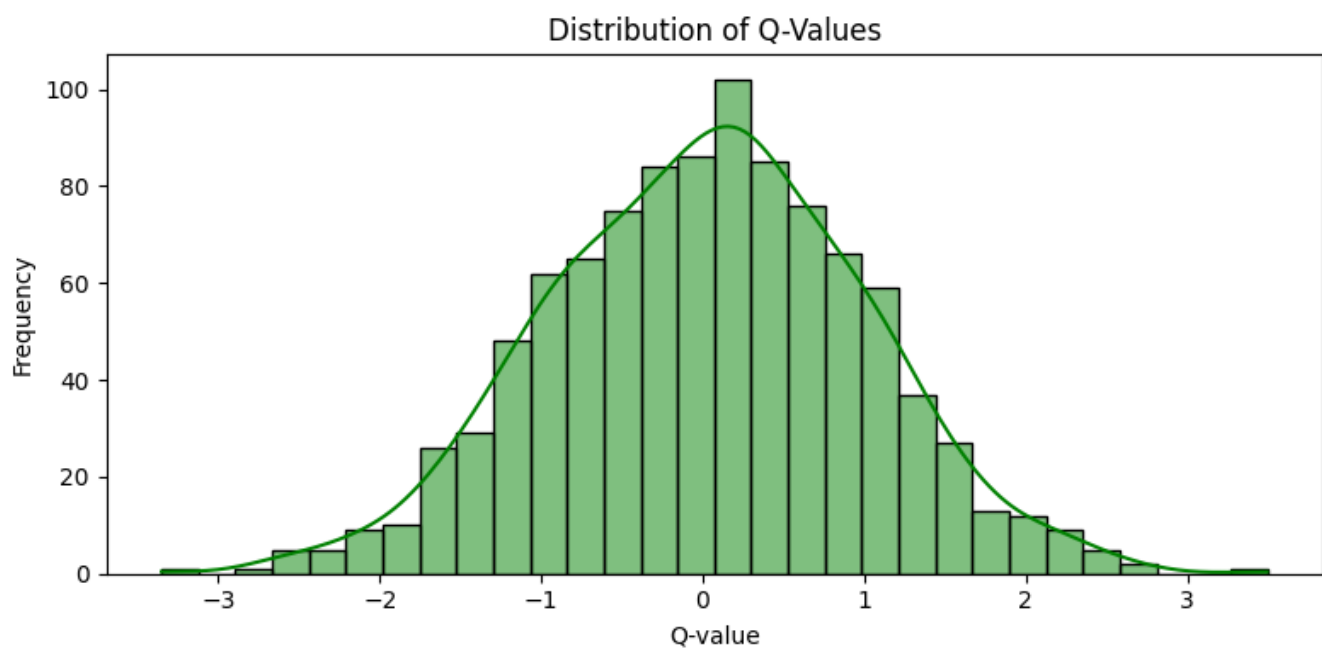


Figure 7: Distribution of Q-values. Reflects the agent's estimation of expected future rewards.

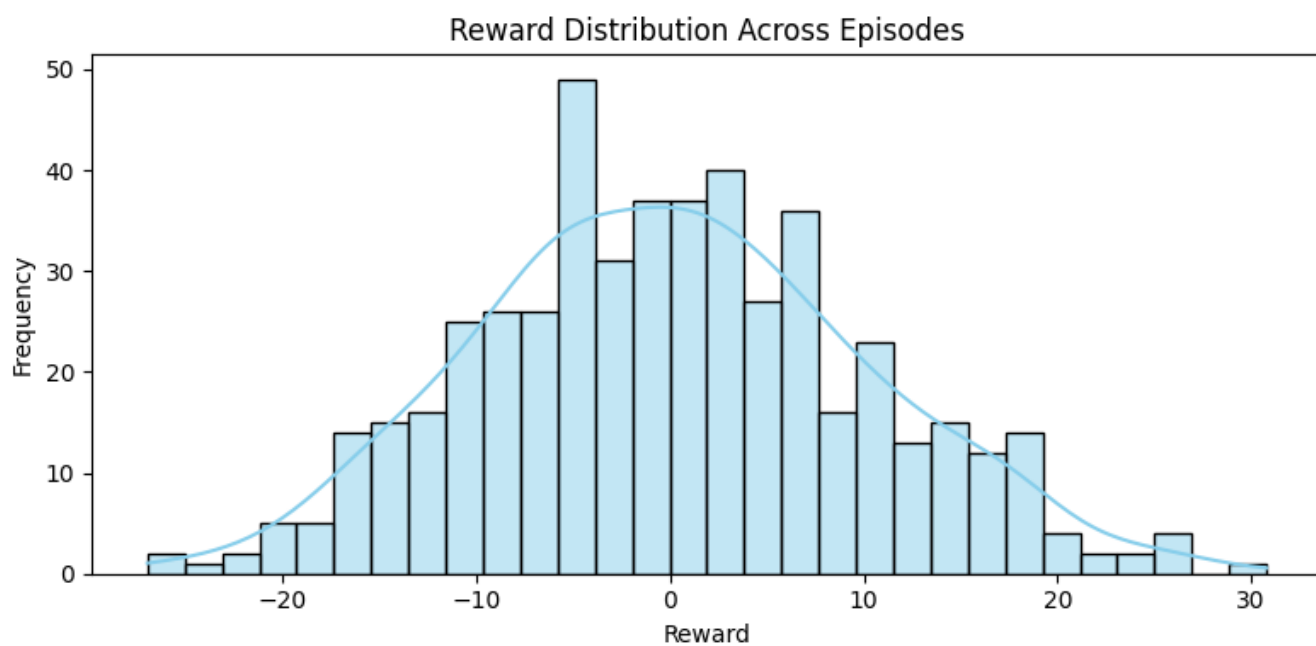


Figure 8: Distribution of rewards per episode. Evaluates consistency and effectiveness of the agent's policy.